

Orcas Infrastructure

Bhavin Gandhi

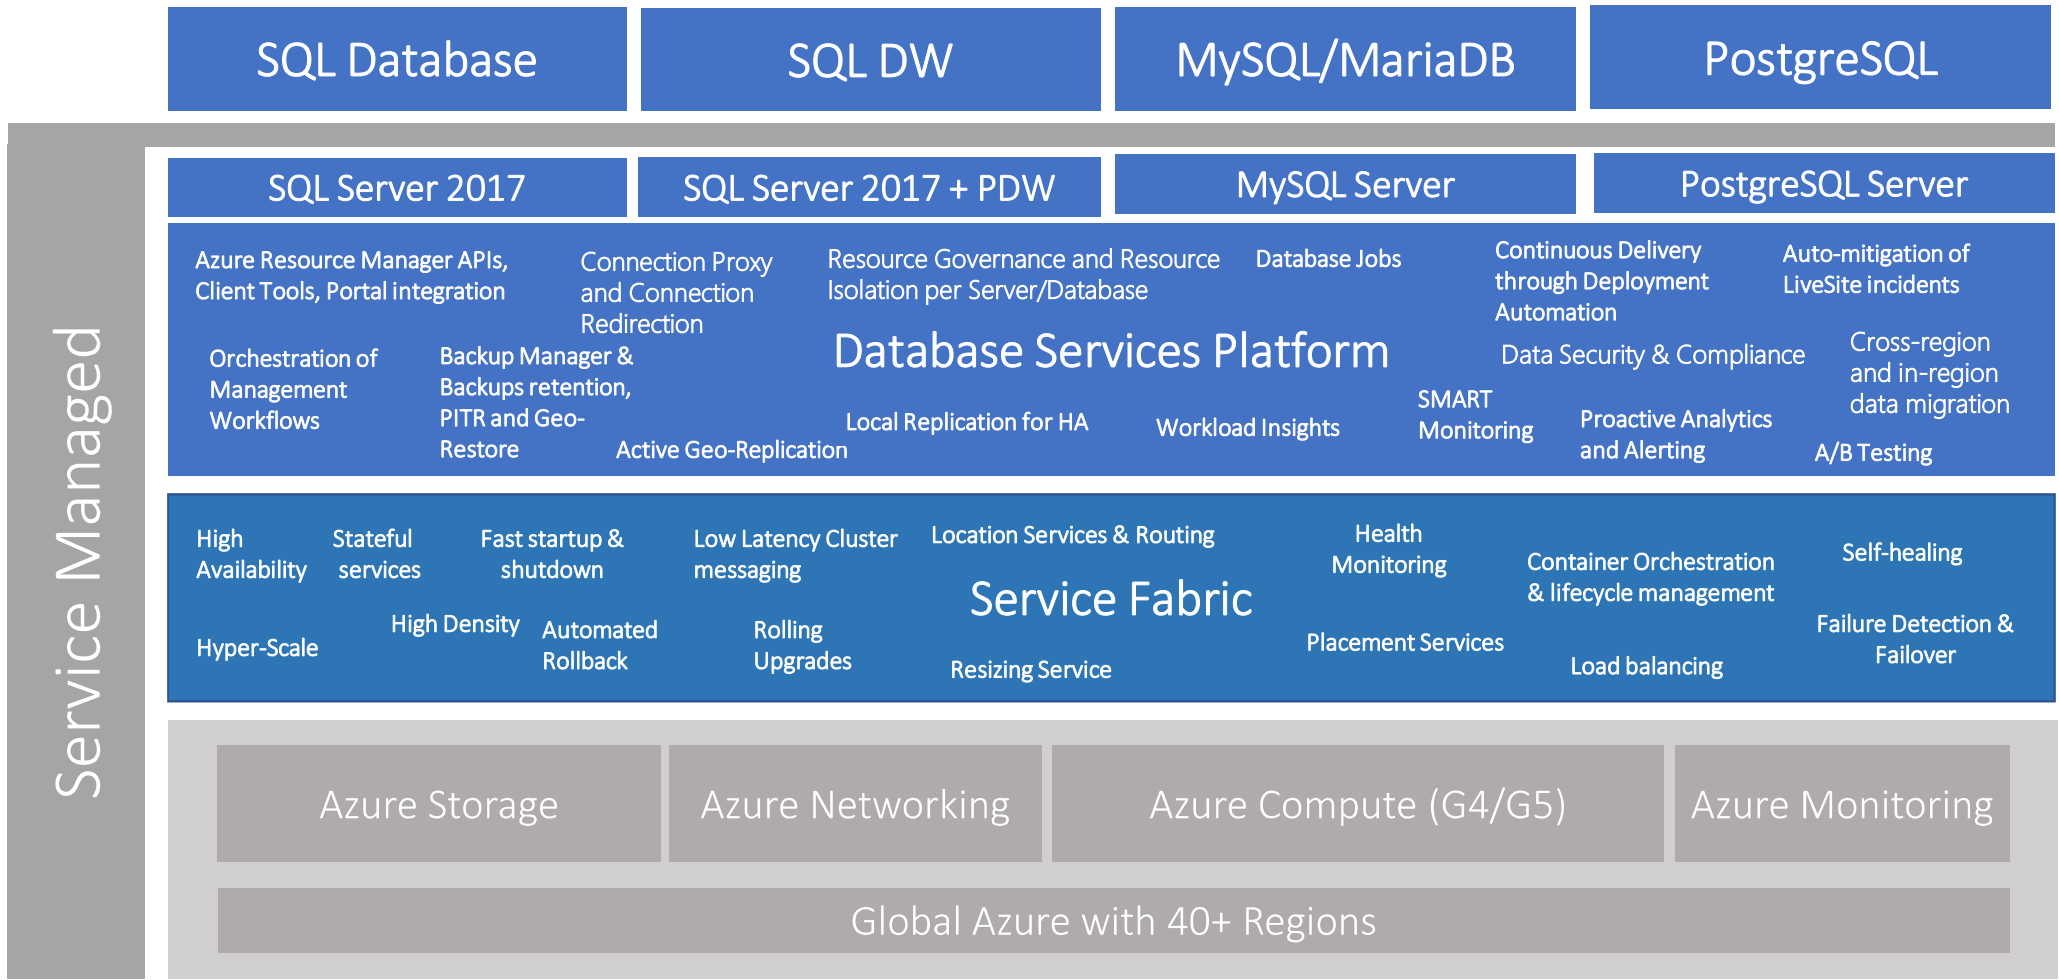
Context

- Relational Database Service in Azure
 - Part of our developer story for bringing customers and their apps on our cloud platform
- Provision and Manage cloud databases for developers/DBAs/CSVs
 - High Availability (HA)
 - Automatic Backups/On-demand Point-In-Time Restores (PITR)
 - Disaster Recovery (DR),
 - Performance tuning/Recommendations
 - Security/Compliance/Auditing
 - On-demand scale up/down
 - Secondary read replicas, etc. etc.
 - **all while saving costs for customers and us**

Context

- Expand our Cloud Database offering to open source database engines
 - Offer choice to our customers for bringing their apps on Azure
- Started as an Incubation project with PostgreSQL as the first PoC and soon MySQL, and now MariaDB and possibly others.
- Leverage existing SQL DB investments (a.k.a Sterling platform) instead of starting from scratch
- Formal name for these services:
 - Azure Database for [PostgreSQL | MySQL | MariaDB]
 - <https://azure.microsoft.com/en-us/status/>


Azure Relational Database Platform



Customer Experience

- Manageability
 - CRUD operations on servers/databases
 - Firewall settings/VNET endpoints
 - Server logs
 - Threat detection
 - Performance Insights
 - Typically done through Azure Portal (portal.azure.com) or Azure CLI
- Connectivity
 - DB Admin tools/data applications connecting to databases to run queries/workload (OLTP)
 - psql/pgAdmin, workbench, or language specific drivers such as libpq, Npgsql ADO.Net driver, ODBC connectors etc.
- **These two distinct experiences are implemented through different services in the architecture stack**

SQL DB Cluster (Region)

**Cluster name:**

wasd-prod-westeurope1-a

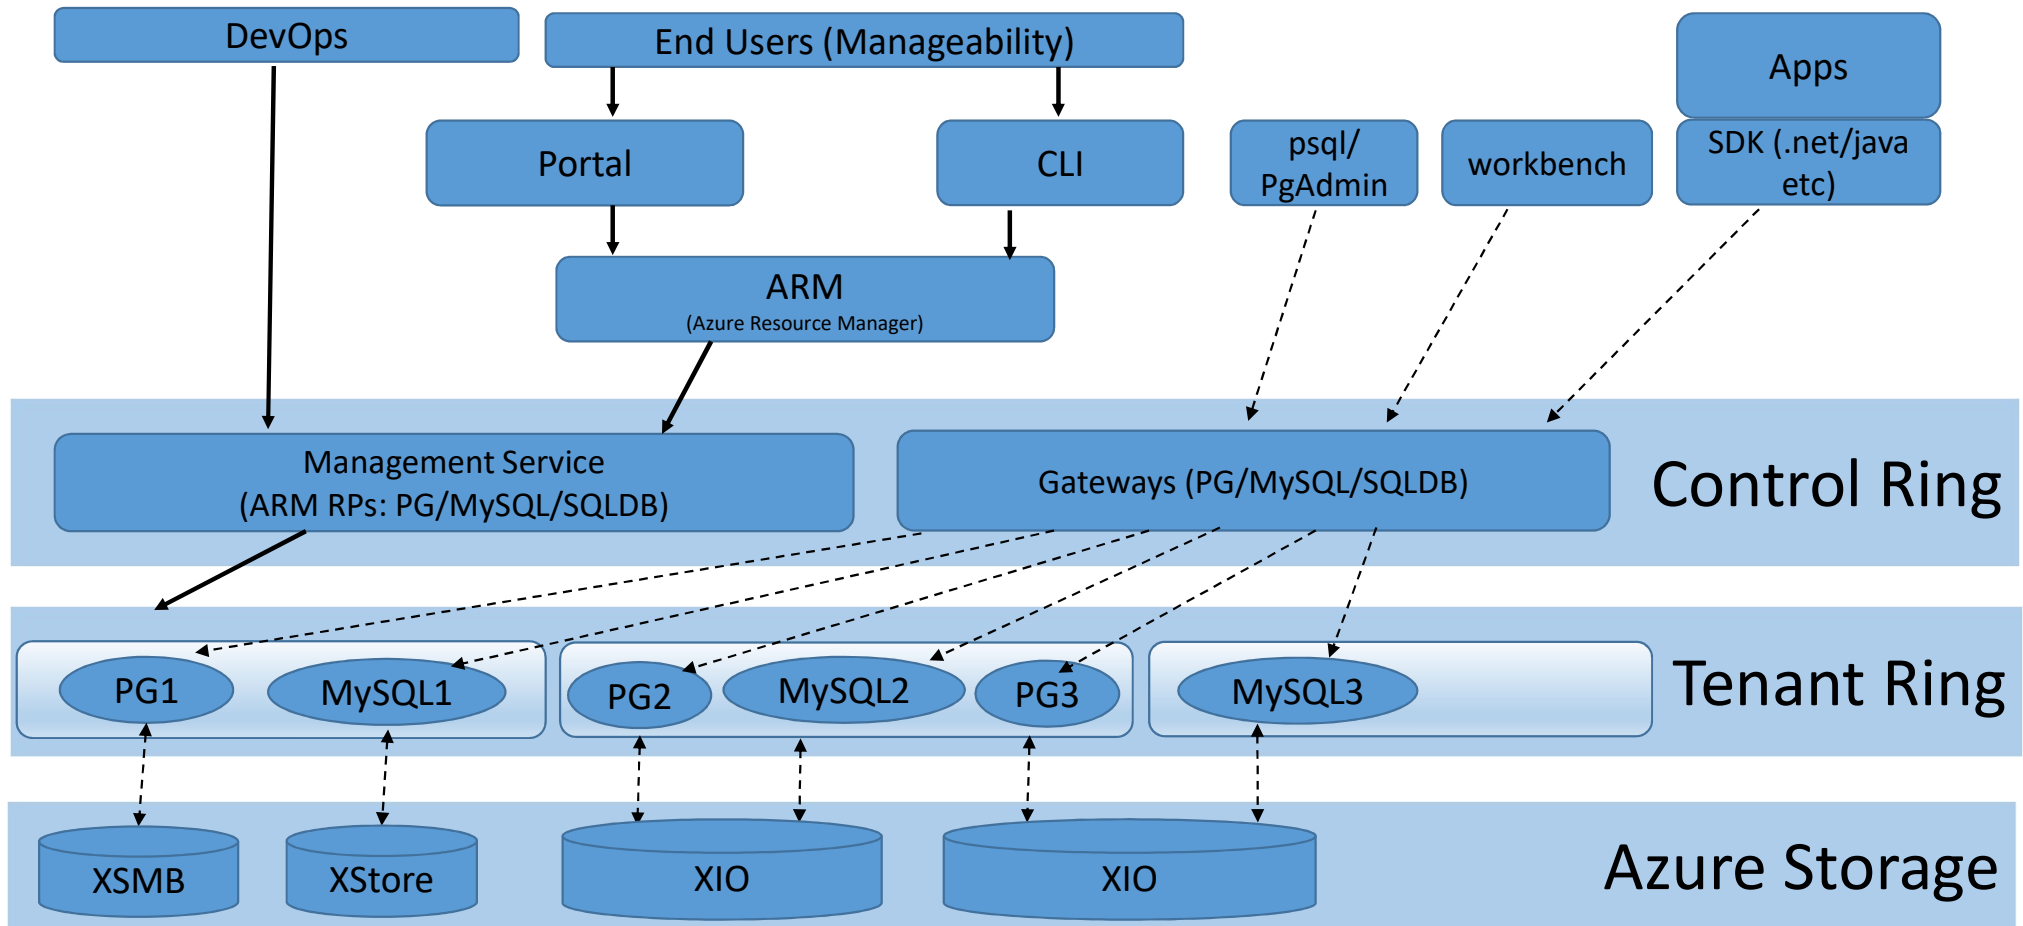
WASD – Windows Azure SQL Database

Prod – Environment (others test, stage)

Region<x> - xth cluster in that geography

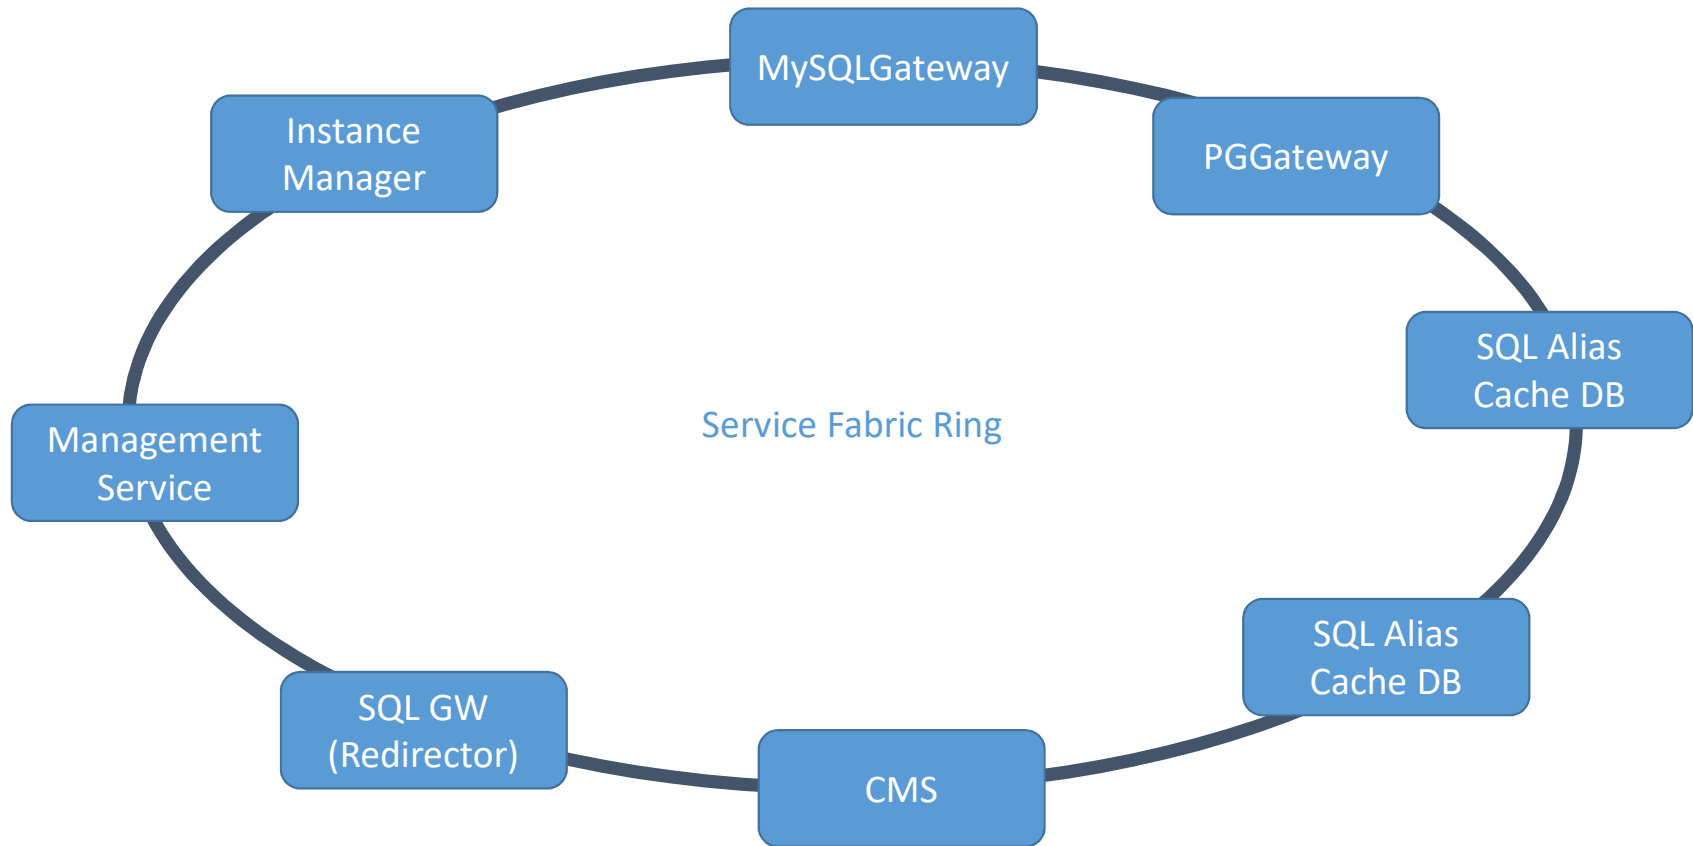
a/b – re-build of a cluster

High level Architecture



SFE View

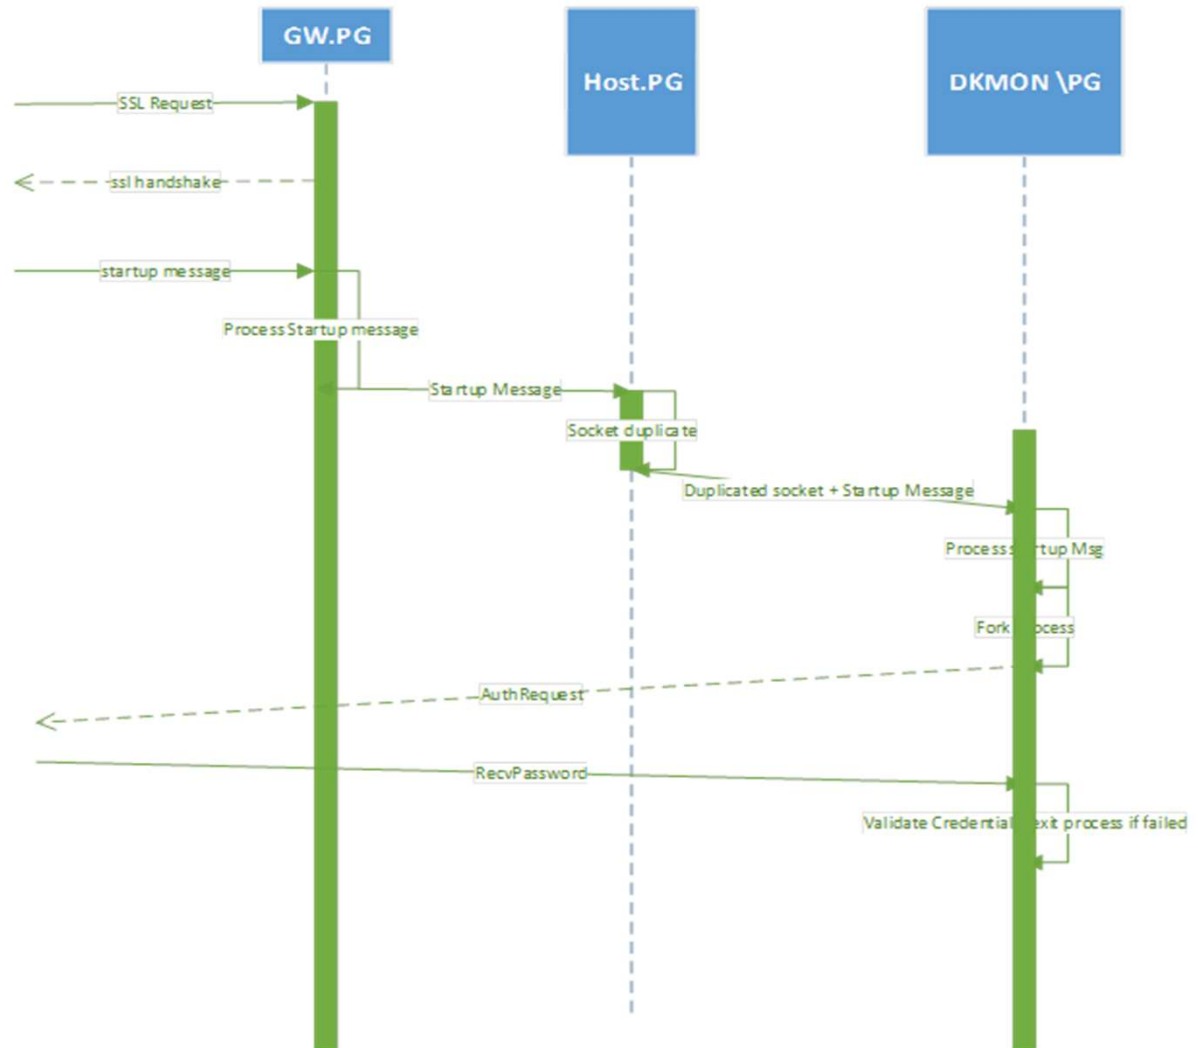
Control Ring Services



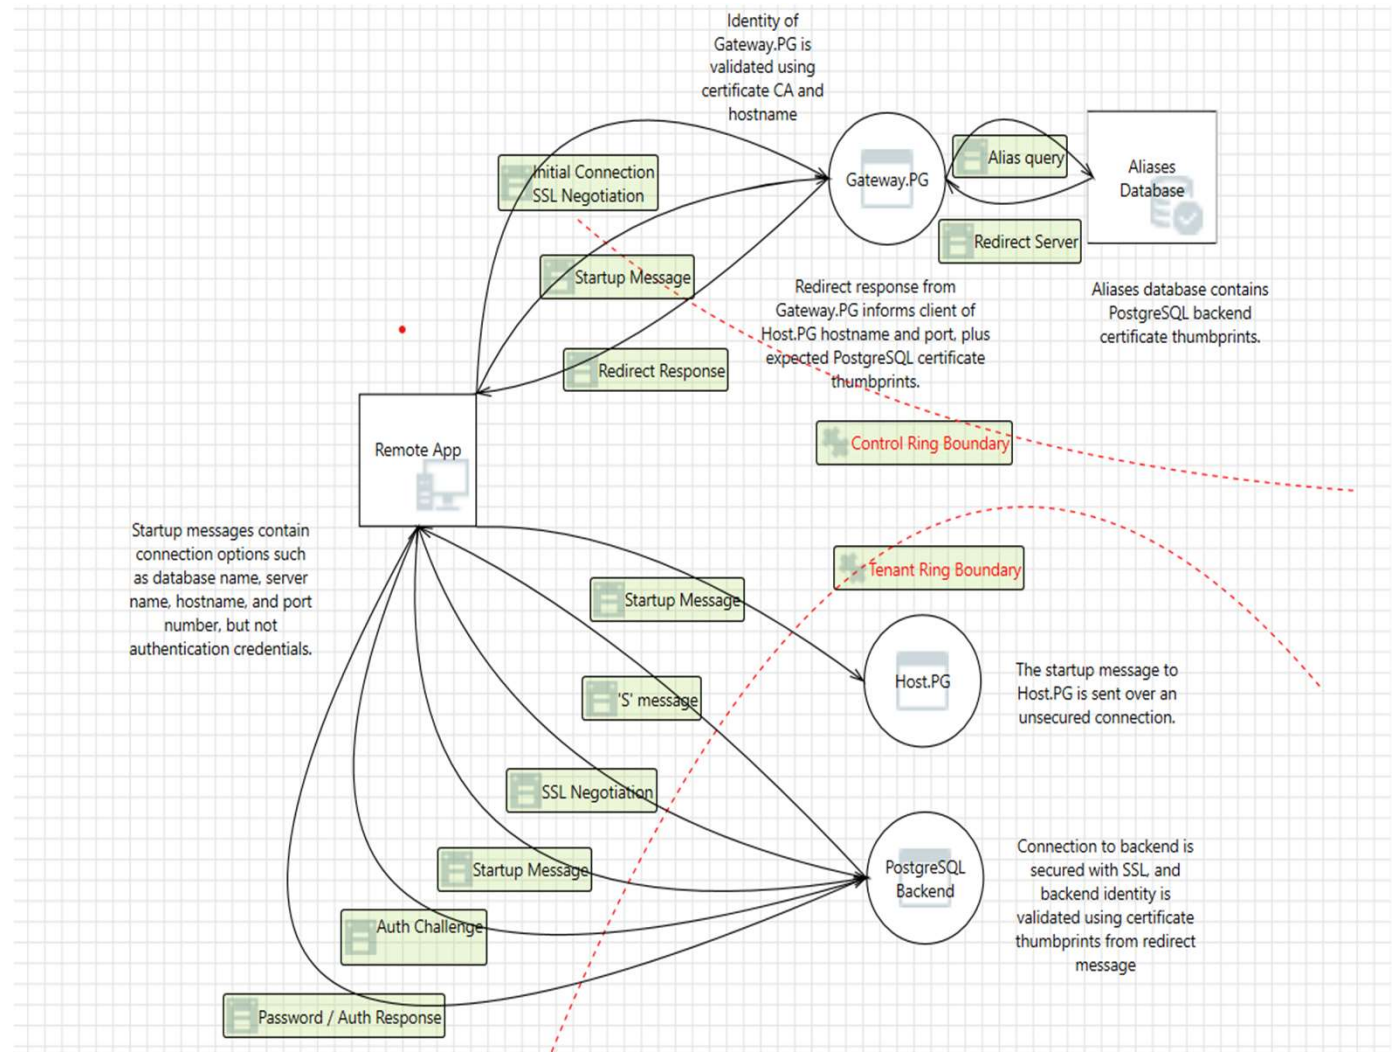
What does the Control Plane Do?

- Handles SQL/PG/MySQL Logins (Gateway services)
- Acts as the resource provider for ARM (ManagementService)
- Knows how to provision SQL Logical servers, PG/MySQL Servers
- CRUD operations
 - **Create:** Create a PG server
 - **Read:** Get the list of servers/databases
 - **Update:** Change the size (SLO/Storage) of this server
 - **Delete:** Delete this server
- Also orchestration of lots and lots of other workflows (PITR, Deployment, Security, Workload Insights, Geo-Restore, Deactivation, etc.)

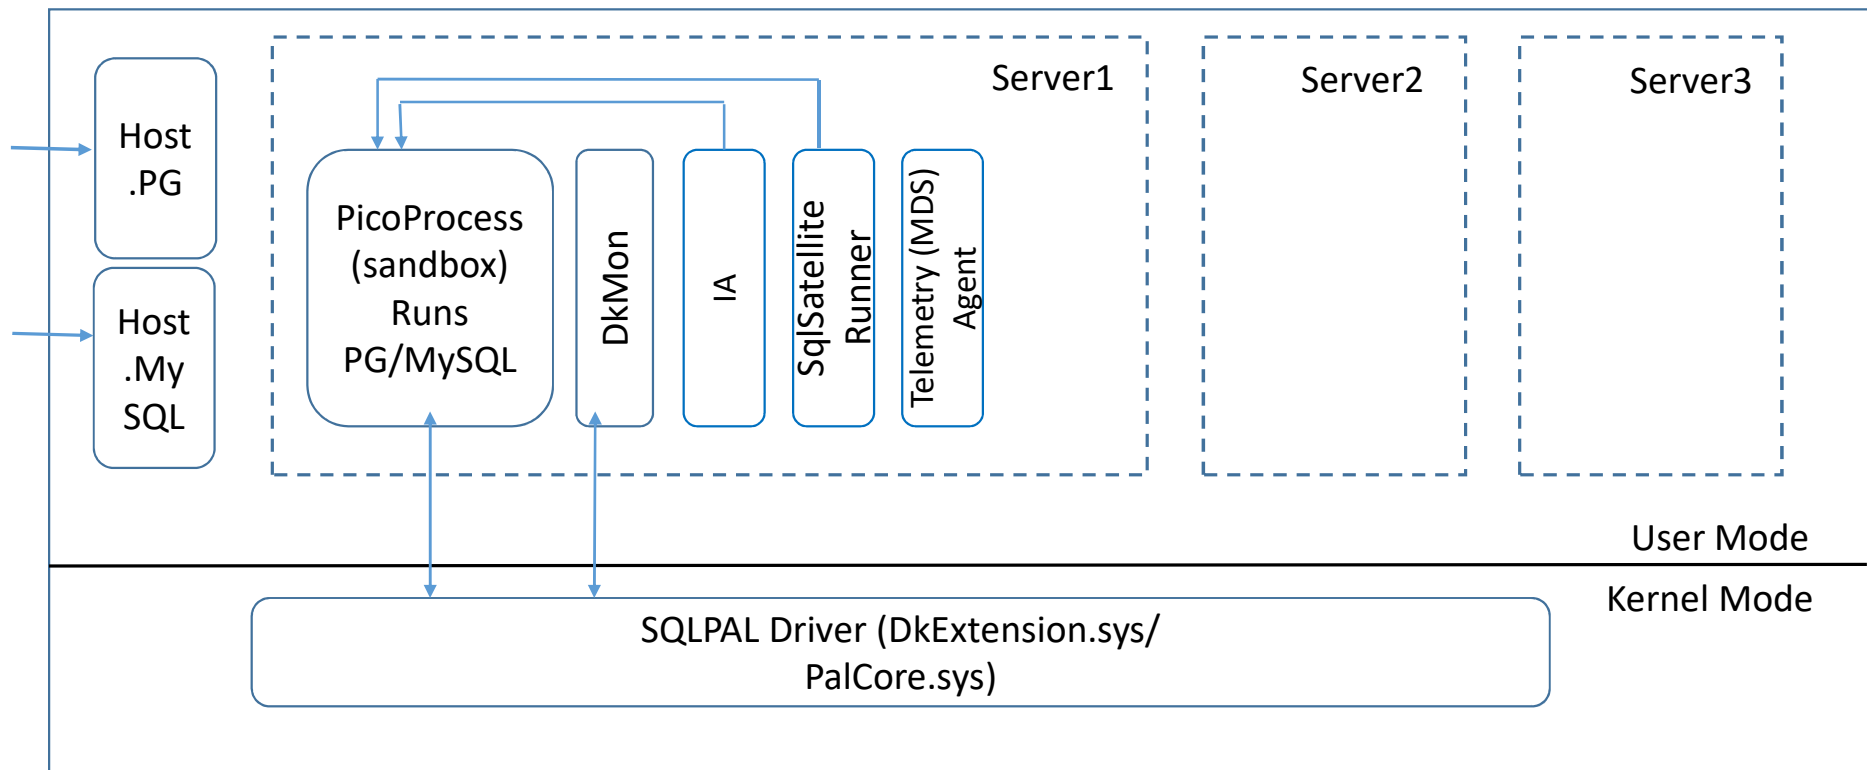
Login Flow (Default)





Login Flow (Redirection)



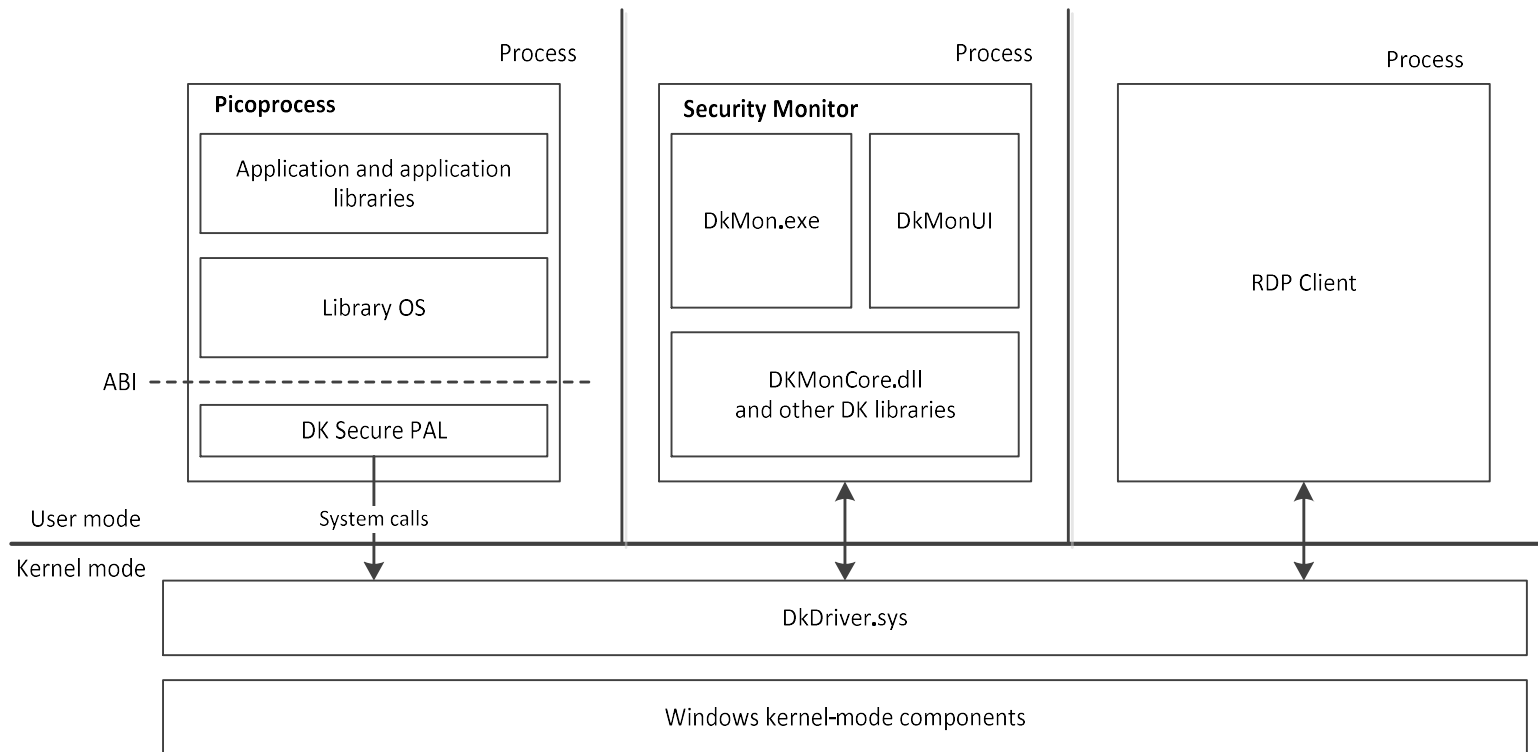
Tenant Ring (DB node view)



SQLPAL

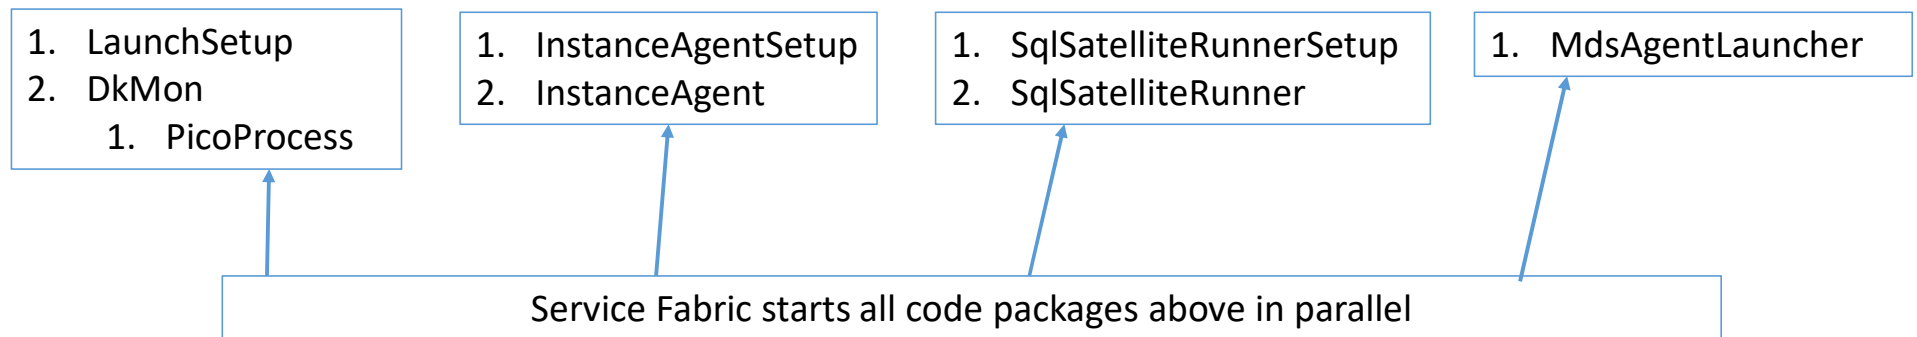


SQLPAL

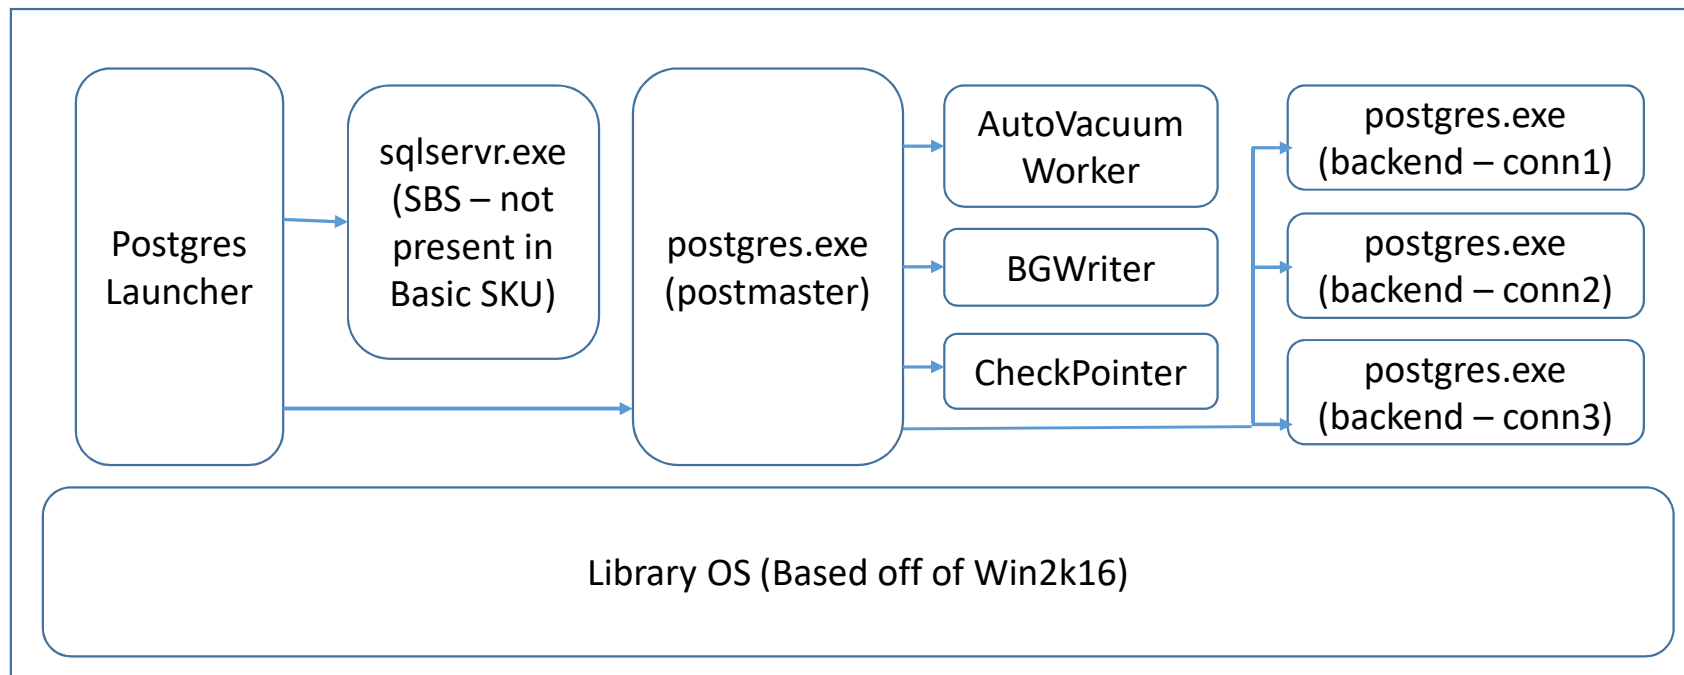


Server startup

- Service fabric orchestrates starting up all services (and restarting them in the event of planned/unplanned failures)
- Below 4 code packages



PicoProcess View (for PostgreSQL)



Useful Telemetry Tables (Kusto)

- MonManagement (For Manageability requests)
- MonLogin (filter by AppTypeName => Gateway.PG/Gateway.MySQL/Host.PG/Host.MySQL)
- MonRdmsPgSqlLaunchSetup (LaunchSetup for PG)
- MonRdmsPgSqlSandbox (DkMon & PG Picoprocess telemetry)
- MonRdmsMySqlLaunchSetup (LaunchSetup for MySQL)
- MonRdmsMySqlSandbox (DkMon & MySQL Picorprocess telemetry)
- MonRdmsInstanceAgent (for IA logs)